

W. Shakespeare

Categories and functors:
much ado about nothing?

Bachelor thesis

26 April 1564

Thesis supervisors: prof.dr. G.A. Bredero
dr. J. van den Vondel



Leiden University
Mathematical Institute

Contents

1	Introduction	3
1.1	Assumptions	3
2	Relevant mathematics	4
2.1	Voronoi	4
2.2	Neighborhood search algorithms	4
2.2.1	Tabu search	5
2.2.2	Simulated annealing	5
3	General model	6
3.1	SAMP	6
3.2	Network structure	6
3.3	Transit Assignment	7
3.4	Costs	7
3.5	Access	8
3.6	The model	9
3.7	General algorithm	9
3.8	Chicago application	9
4	Case Study: Leiden	11
4.1	Transit system	11
4.2	Facility and Community nodes	12
4.3	Travel demand	12
4.4	Chosen parameters	12
5	Explanation of experiment	14
6	Results	15
7	Discussion	16
A	Pseudocode	17
B	Data used	20
B.1	Communities	20
B.1.1	PC4	20
B.1.2	Voronoi	20
B.2	Travel data	21

Chapter 1

Introduction

1.1 Assumptions

The following things are assumed to be true

1. The group of people that the model is made for is very small compared to daily travelers, therefore this small group will have no effect on overcrowdedness or other usage related metrics.

Chapter 2

Relevant mathematics

2.1 Voronoi

Following the definition in [Møl94], a voronoi diagram is given by a set $\Phi = \{x_i\}$ of weightpoints in a metric space X . Each weightpoint generates a region

$$C(x_i|\Phi) = \{y \in X : d(x_i, y) \leq d(x_j, y) \ \forall x_j \in \Phi\},$$

where d denotes some distance function. It follows that each voronoi region $C(x_i|\Phi)$ contains all points that have x_i as nearest weightpoint.

Applications of voronoi diagrams can be found in many fields. Consider for example the placement of automated external defibrillator (AED) boxes in a city. An even spread of these boxes would theoretically save the most people. The positions of these boxes could be used as weight-points, with the euclidian distance in $\mathbb{R}^2$. This results in a voronoi diagram where smaller regions represent boxes that are relatively close to each other, while larger regions represent boxes that are far from others. This information could be used to move or add boxes, and thus getting a more even distribution.




Figure 2.1: Entrance of Gorlaeus building of Leiden University, with voronoi diagram as decoration. Picture taken from [Swe25].

2.2 Neighborhood search algorithms

When an optimization problem has a large but finite solution space, a method of traversing possible solutions is a neighborhood search (NS) procedure. Given a solution space S , a neighborhood function $N : S \rightarrow \mathcal{P}$ and a feasible solution $s \in S$, a neighborhood $N(s)$ of s can be found. This contains all feasible solutions that can be reached from s after doing one *move*.

General NS algorithms are heuristic approaches and use a starting solution s^* and try to find a solution in $N(s^*)$ that has a better value than s^* . If such solution s exists, a new iteration is done where a more optimal solution is sought in $N(s)$. This iteration loops until a solution does not have a more optimal solution in its neighborhood.

NS algorithms have no method of leaving local optima, and thus get stuck often. To prevent this behaviour, many methods have been developed. We discuss two of those methods, both relying on making suboptimal moves. Just like NS, these are heuristic approaches to optimization problems.

2.2.1 Tabu search

Tabu Search (TS) is an NS algorithm that uses *Tabu lists* to escape local maxima. Visited solutions or executed moves are stored in *short-term memory* (STM) and can not be revisited or redone for a small number of iterations, this is to prevent cycling. The length of time that solutions/moves are stored in STM is dependent on *Tabu tenures*. *Long-term memory* (LTM) is used to store solutions with a good objective value, so they, and their neighbors, can be revisited after the algorithm becomes trapped.

Tenures can be adjusted dynamically while running the algorithm. This is to allow a greater focus on *diversification*, the exploring of many different solutions, or on *intensification*, the exploiting of good solutions.

2.2.2 Simulated annealing

Simulated Annealing (SA) is named after a technique in metallurgy where a material is heated and cooled in a controlled fashion, resulting in a change in physical properties. SA uses the same iteration as a basic NS algorithm, comparing different neighbors to choose a good next move. When choosing a next step, a temperature parameter T is used to randomly choose a worse option. This parameter decreases while running, resulting in erratic moves in the beginning and more controlled moves towards the end of running.

Chapter 3

General model

This section summarizes the general model and algorithm developed by Rumpf and Kaul in [RK21]. As the model and algorithm is made by them, a more detailed explanation can be found in their paper. In this section we will give a summary so the later parts of this study can be understood.

3.1 SAMP

The algorithm made by Rumpf and Kaul in [RK21] was made to solve the *social access maximization problem* (SAMP). The formulation of the problem is very general as to be applicable to many situations.

Let L be the set of considered transit lines, where a transit line $l \in L$ is a series of stops with a given travel time between each stop. A schedule y is a vector with $y_l \in \mathbb{N}$ the amount of vehicles assigned to line l in a given time period. The vehicles are distributed evenly over the time period, so 5 vehicles in an hour gives 1 trip every 12 minutes.

The initial schedule is denoted as y^* . The model has upper and lower bounds $y_l^{\min} \leq y_l \leq y_l^{\max}$ for all $l \in L$. These can be set to 0 and infinity if no restriction is needed, or both to y_l^* if no change is wanted.

Let x be the user flow vector for the transit network, x^* the initial flow vector. User flow represents how busy parts of the network are, which is influenced by the schedule y and found by the $\text{TransitAssignment}(y)$ function, which will be explained further in Section 3.3.

Other restraints of the model are given by operator and user costs with upper bounds $B_{operator}$, B_{user} respectively. Operating costs are mostly made up by monetary costs of running the transit network and monetary gain of users paying to use the service. User costs are made up of difficulty in traversing the network, be that in monetary cost or time spent in the system. More on these costs in Section 3.4.

Let $\text{Access}(y)$ be a function describing the social access objective. This value only depends on the design of the network y and not on user flow x , because the model is designed for a small group of individuals following Assumption 1 Different choices for $\text{Access}(y)$ are given in Section 3.5.

3.2 Network structure

The model uses a representation of the transit network to calculate traveltime and routechoices. For this, a directed graph $G = (V, A)$ is used. V is the set of all relevant nodes, be that origin, destination, transit stop nodes. A is the set of arrows linking the nodes together, these represent different modes of travel like walking, bus or trainriding.

The node set is partitioned in subsets, each containing specific nodes. Let $V_{stop} \subset V$ contain nodes representing the physical location of transitstops. Let $V_{board} \subset V$ be the set of boarding nodes. Each stopnode $i \in V_{stop}$ has boardingnodes $j \in V_{board}$ for every line $l \in L$ that stops at that stop. Boarding nodes are used to distinguish between transit lines. Let $V^+, V^- \subset V$ be origin- and destination nodes, respectively. These represent the physical places where trips begin and end. To calculate social access, subsets $\tilde{V}^+, \tilde{V}^- \subset V$ are added, representing community and facility nodes, respectively.




Figure 3.1: Example of network representation containing two lines and three stops. Each stop has nodes in V_{stop} . Line 1 stops at all stops and thus has three boarding nodes, line 2 only has boarding nodes at stop A and stop C. Boarding nodes are connected to their stops with boarding arcs and alighting arcs. Boarding nodes are connected by line arcs if a transit line connects their respective stops. Stop A and stop B are connected by walking arc, because they are close enough. Figure taken from [RK21].

Just like the node set, the arc set A is partitioned in subsets. Let $A_{line} \subset A$ be the arcs that represent the journey made by transit lines and let $A_{board}, A_{alight} \subset A$ be the arcs representing entering or leaving a transit vehicle at a certain stop. Let $A_{walking} \subset A$ be walking arcs, connecting different nodes if their physical distance is walkable.

3.3 Transit Assignment

The transit assignment model, $\text{TransitAssignment}(y)$, models the way daily users interact with a given schedule y . It assumes users try to take their optimal route, which may not be the fastest due to overcrowdedness. The flow produced by the model is called *user-optimal flow* and is often less optimal than when all flow is controlled directly.

There are many such models, but Rumpf and Kaul choose the nonlinear model by Spiess and Florian in [SF89]. This model increases arc cost based on crowdedness, discouraging flows from exceeding line capacity. The user-optimal flow output can not be changed in a way that individual users benefit from changes in only their route. The model is relatively simple, a desired property, as the model is run many times during the algorithm.

3.4 Costs

Generally, operator costs are monetary costs relevant to running the transit system. Things like purchasing and maintaining vehicles, employing drivers and other staff and other costs. In this study, we restrict the total number of busses to not increase. This gives no reason to purchase new vehicles or employ more drivers, so we assume that operator costs will not increase with a new schedule. This means that we can omit the operator cost constraint from the model.

User costs are time costs made by using the system. Things like travel time, wait time, number of transfers and discomfort. All the user costs can be weighted based on transfer mode, 5 minutes of walking may weigh more than 5 minutes of busriding for example. For each $(i, j) \in A$ a flow amount x_{ij} and cost c_{ij} can be implemented. Flow amount represents the amount of users taking that arc, while cost represents some penalty in using that arc, like time needed to traverse it. The user cost function takes the form

$$\text{UserCost}(x, \omega) = \theta_1 \sum_{ij \in A_{line}} c_{ij} x_{ij} + \theta_2 \sum_{ij \in A_{walk}} c_{ij} x_{ij} + \theta_3 \omega,$$

where $\omega \in \mathbb{R}$ is the total time spent waiting at busstops by all users combined.

Due to Assumption 1 we do not want the total user cost to increase too much, so we introduce a parameter ε to give a constant upper bound

$$\text{UserCost}(x, \omega) \leq (1 + \varepsilon) \text{UserCost}(x^*, \omega^*).$$

When choosing ε , a balance needs to be found between a large increase, giving a larger set of possible solutions and possibly a more optimal one at the cost of daily user experience.

3.5 Access

The objective function $\text{Access}(y)$ should describe how easy it is to travel from a general community to a general facility, with a focus on equity of access. To do this, a metric is made to measure the accessibility for a specific community $i \in \tilde{V}^+$.

The chosen metric $A_i(y)$ is *gravity-based*, meaning that the attractiveness of facilities reduces gradually as distance increases. An alternative is a discrete distance based approach, where facilities $j \in \tilde{V}^-$ are considered only if they are within a certain distance from the source community i . This approach is not preferred, as facilities that lie just inside the considered radius are weighted equally to facilities that are nextdoor, and facilities that lie just a bit further are not considered at all. This is solved by including a distance function $d_{ij}(y)$, measuring traveltime within G from i to j , and a gravity decay parameter $\beta > 0$. A larger β will result in a higher penalty for further travel.

As the users can choose between facilities, a competition-related metric $F_j(y)$ is used to find the crowdedness of facilities. It is defined as

$$F_j(y) = \sum_{k \in \tilde{V}^+} P_k d_{kj}^{-\beta}(y) \quad \forall j \in \tilde{V}^-$$

where P_k is the population of community k and $d_{kj}^{-\beta}(y)$ is the previously mentioned gravity decay. This crowdedness can be compared to the quality S_j of a facility, usually taken as some metric of patient capacity.

The accessibility metric for a given region is then defined as

$$A_i(y) = \sum_{j \in \tilde{V}^-} \frac{S_j d_{ij}^{-\beta}(y)}{F_j(y)} \quad \forall i \in \tilde{V}^+.$$

This metric was first developed in [Wei76] to describe differential access to employment opportunities, but can also be used for other competition based access cases.

Since the goal is to create more equity in accessibility, not all communities should be included in the overall objective function. Instead, the communities with the worst accessibility should be improved. By maximizing the minimum communities, a more equitable scenario is achieved. The model parameter $\mathcal{K}$ is introduced to achieve this. The objective function is then defined as

$$\text{Access}(y) = \sum_{\substack{\mathcal{K} \text{ minimum elements of} \\ \{A_i(y) | i \in \tilde{V}^+\}}} A_i(y).$$

Note that the objective function will be evaluated many times, and that the $\mathcal{K}$ worst communities are dependent on the y of the current iteration, so the evaluated and improved communities changes while the algorithm is running.

The parameter $\mathcal{K}$ needs to be chosen between 1 and $|\tilde{V}^+|$, where extreme values are not beneficial to the outcome. Low values for $\mathcal{K}$, like 1, may result in minor improvements to a single community with no regard for the relatively large effect on other communities. High values, like $|\tilde{V}^+|$, may result in making large improvements to the highly accessible communities at a cost of making the less accessible communities even worse, as long as the total accessibility increases.

3.6 The model

The SAMP model takes the shape of the following nonlinear program

$$\max_y \text{Access}(y) \quad (3.1)$$

$$\text{s.t.} \quad \sum_{l \in L} y_l \leq \sum_{l \in L} y_l^* \quad (3.2)$$

$$y_l^{\min} \leq y_l \leq y_l^{\max} \quad \forall l \in L \quad (3.3)$$

$$y_l \in \mathbb{N} \quad \forall l \in L \quad (3.4)$$

$$x = \text{TransitAssignment}(y) \quad (3.5)$$

$$\text{UserCost}(x, \omega) \leq (1 + \varepsilon) \text{UserCost}(x^*, \omega^*) \quad (3.6)$$

Where Eq. (3.1) is the objective function described in section 3.5, Eq. (3.2) constrains the total number of vehicles to not increase from the initial amount, Eq. (3.3) ensures upper and lower bounds for vehicle allocation, Eq. (3.4) ensures a whole number of vehicles is allocated to each line, Eq. (3.5) defines user flow as discussed in section 3.3 and Eq. (3.6) gives an upper bound to the increase in user cost.

3.7 General algorithm

The SAMP model can be solved heuristically with the following algorithm. It is a combination of *Tabu search* (TS) and *Simulated annealing* (SA), two versions of neighborhood search mentioned in Sections 2.2.1 and 2.2.2. Full pseudocode can be found in Appendix A.

Using the current transit schedule as a starting point, a neighborhood search is conducted, where neighbors of a solution s are defined as valid solutions that differ only from s by removing one vehicle from a transit line, and adding that vehicle to another line.

During each iteration, the two best neighbors are determined. The best neighbor is chosen as the solution for the next iteration, unless the SA process forces the second best option to be chosen. The option that is not chosen gets put in *long term memory* to be revisited later, while the adding/removing vehicle moves are put in *short term memory* to not be repeated soon, according to TS.

The chance that suboptimal neighbors are chosen is determined by a temperature parameter T , that decreases while the algorithm is running. This ensures more erratic exploration at the start of the algorithm and more exploitation of good solutions towards the end.

After a predetermined amount of iterations is completed, all solutions in the long term memory and their neighbors are revisited. The most optimal solution that is found is then locally optimal.

3.8 Chicago application

In their paper, Rumpf and Kaul applied their model to the transit system of Chicago [RK21]. In their application, a transit stop set V_{stop} with 997 different stops was used. These stops were also used as origin-destination (OD) pairs for day-to-day travel volume. Community nodes $\tilde{V}^+$ were defined using the $|\tilde{V}^+| = 77$ Chicago community areas, where node position was taken as population-weighted centroids, based on 2010 census data. Facility nodes $\tilde{V}^-$ were the primary care facilities, of which $|\tilde{V}^-| = 119$ are within the Chicago area. Each facility was given a quality of $S_j = 1$, $\forall j \in \tilde{V}^-$. Walking arcs were created when different locations were 0.75 miles (approximately 1.2 km) within each other, walking speed was set at 4 ft/sec (approximately 4.4 km/h).

Transit lines were acquired from General Transit Feed Specification (GTFS) data. It includes 126 bus lines and 8 train lines. It was also used to estimate the current number of vehicles y_l^* for each line $l \in L$. The time horizon τ was taken as 24-hour period. It was assumed a single type of bus was used, with capacity $b_l = 39$ for all buslines, and a single type of train with capacity $b_l = 228$ for all trainlines. Fleet bounds $y_l^{\min}, y_l^{\max}$ were set to equal y_l^* for all train lines to force these to remain the same. For bus lines $y_l^{\min}$ was set to an amount to achieve at least one arrival per 30 minutes, no maximum $y_l^{\max}$ was set.

Their computational trials also used express lines, existing bus lines only servicing approximately 20% of their initial stops. With less stops, a lower circuit completion time τ^l could be

achieved, thus increasing their effective frequency f_l . The stops that remained were in highly accessible communities or in badly accessible communities, so the lines could connect badly accessible communities to facilities. These express lines had initial fleet sizes of $y_l^* = 0$.

The algorithm was run for 500 iterations, followed by a neighborhood search of the best found solution to ensure local optimality. The access function in Section 3.5 was evaluated with gravitational decay parameter $\beta = 1.0$ and community count of $\mathcal{K} = 8$ (approximately 10.39% of 77 communities). The user cost function in Section 3.4 used equal mode weights $\theta_1 = \theta_2 = \theta_3 = 1$ and maximum relative increase $\varepsilon = 0.01$.

They also ran trials where either user cost constraints were ignored, only allowing vehicle swaps between lines and their express lines, and excluding express lines.

Chapter 4

Case Study: Leiden

This section will discuss our implementation of the same model on Leiden. It will also go over the different choices of parameters and motivate these choices.

The municipality of Leiden is much smaller than the Chicago area. This naturally gives a lower amount of busstops, buslines, health facilities and so forth. Although there are multiple train stations in Leiden, the trainsystem is not taken into account for this study.

4.1 Transit system

As Qbuzz, the company running the bussystem in Leiden, operates in a greater region, most of their buslines lie partially or completely outside of the study area. For this study, only buslines that travel at least partially through the study area are considered. A busline usually drives a linear route between two points, stopping at bus stops in between. For the purposes of this study, the back and forth trip is considered a complete circuit.

Buslines that service busstops outside the study area are modeled to have all stops within the area and a single stop at their turnaround point. For example, line 20 between Leiden Centraal and Noordwijk Vuurtorenplein services 26 stops in total, 4 of which are within the study area. It is modeled to have stop nine times per circuit, twice per stop within the municipality and once at Vuurtorenplein. This way of modeling ensures the circuit time τ_l works correctly for each line l .

One notable exception to this method are the linking of lines 1, 3 and 4. Their scheduled routes end where the next route begins. Line 1 travels between De Vink and Leyhof, line 3 between Leyhof and Merenwijk and line 4 between Merenwijk and De Vink. The vehicles allocated to these lines currently continue the next line after completing their given trip. We will keep this planning choice and thus model these lines as two lines, one for each direction.

Since Leiden is much smaller than the Chicago area, the implementation of express lines was excluded.

After set of buslines L was selected, a schedule was retrieved from [Qbu25]. The chosen schedule was for monday the 10th of february of 2025, between 14:00 and 15:00 in the afternoon. This weekday and timeframe was chosen as a ordinary, non-rush hour time, that represents ordinary bus travel.

The information given by a schedule is average traveltime between stops, wait time at a stop and amount of busses per line per hour. The used algorithm takes busses per line per time-horizon τ and averages this out to arrive at the frequency of a line, where more busses allocated over a time-horizon would mean a higher frequency of service. The circuit time τ_l of a line l is the time needed to complete a round trip, after which it can be used to complete a new trip. Buslines with a lower circuit time can thus get a appropriate frequency while having a lower amount of busses allocated to them, relative to lines that have longer circuit times.

Names of- and traveltime between busstops can be retrieved from the schedule, giving a transit-stop set V_{stop} of 121 different busstops, of which 100 are within the study area. As most bus stops consist of a pair of physical busstops, one for each direction, the midpoint was taken to represent the physical pair in the graph. Distances to other points, like other busstops or community centers, can be determined and walking arcs can be connected accordingly.

For simplicity it is assumed that the operator uses a single type of bus, the Iveco Crossway EV 13m, with a capacity of $b_l = 50$ for all lines $l \in L$.

4.2 Facility and Community nodes

Facility nodes $\tilde{V}^-$ are defined by the general practitioners (GP) offices of Leiden. Taken from [Ned25] a total amount of $|\tilde{V}^-| = 37$ are used. The amount of full-time employed doctors that work at a facility j is taken as the quality S_j of that facility.

For this study, we compare two different methods of defining community nodes $\tilde{V}^+$. The first method concerns itself with postcoderegions in the Netherlands. Each address in the Netherlands has a postcode, consisting of four numbers and two letters. If one only uses the four numbers, the Netherlands can be divided into pc4 regions. There are $|\tilde{V}_{\text{pc4}}^+| = 16$ different pc4 regions within the municipality. Center node location was taken as their middlepoint and each node was assigned population data taken from [All25].

These pc4 regions are still quite large and all, non-day to day trips are simulated to start at the center node to then walk to a bus stop. This may imply that busstops far from the middlepoint of a region are used less then in reality. To remedy this, another way of defining community nodes is found.

According to [Wan22], travelers prefer to take a bus at the nearest bus stop, even when that involves a transfer instead of a direct route. To apply this effect on our study, the regions we choose should reflect trips that start at the closest bus stop. Voronoi diagrams, as defined in Section 2.1, are the ideal method of dividing a plane in sections where each section contains one weightpoint and all points that have that weightpoint as the nearest weightpoint.

Using this method, with busstops as weightpoints, we can divide the municipality in $|\tilde{V}_{\text{vor}}^-| = 100$ regions. We have one region less than the number of busstops, because one busstop lies outside the municipality. We estimate the population data for the voronoi regions by looking at the overlap between each voronoi region and the pc4 regions, and assuming the known population of a pc4 region is distributed evenly over its surface. Let

$$\begin{aligned} \tilde{V}_{\text{vor}}^- & \text{ set of voronoi regions,} \\ \tilde{V}_{\text{pc4}}^+ & \text{ set of postcode regions,} \\ s_v & \text{ surface area for } v \in \tilde{V}_{\text{vor}}^-, \\ d_p & \text{ population density for } p \in \tilde{V}_{\text{pc4}}^+, \\ f(v, p) & \text{ fraction of } v \text{ that lies in } p. \end{aligned}$$

Then, the population of a voronoi region v is given by

$$s_v \cdot \sum_{p \in \tilde{V}_{\text{pc4}}^+} d_p f(v, p).$$

4.3 Travel demand

No actual travel data could be aquired, so this had to be estimated. As the schedule gives the frequencies of service, it also gives the total supply of bus transit. Assuming the supply is an accurate reflection of the demand of bus transit, all busses are roughly equally filled.

We therefore assume that each bus is halfull. The algorithm wants travel demand between busstops, so this data was made by counting the amount of lines servicing a pair of bus-stops directly and multiplying this by the lines' frequency and the capacity of a bus. The resulting table can be found in Appendix B.2.

4.4 Chosen parameters

Since our study tries to improve accessibility for less mobile people, we want to limit the distances walked. Therefore the choice was made to connect nodes via walking arcs if they are within 250 meters of eachother, and walking speed was set at 4km/h. Since all day-to-day travel is simulated to begin and end at a busstop, this change only affects those travelers when walking between busstops for a transfer. It will have a much larger effect the group of less mobile people that the algorithm is made to optimize for.

The time horizon τ was set at 540 minutes, the tipical time for a GP to be open during a weekday. The schedule taken is one based on 60 minutes outside of rush-hour, and is thus

extrapolated to count for the whole day. This is slightly inaccurate, as more busses are used during rush-hour to account for the higher demand. Since the algorithm averages out both the travel-demand and the busses used, accounting for this inaccuracy would be undone by the algorithm.

Fleet bounds are set identically as in [RK21], that being no upper bound $y_l^{\max}$ and lower bound $y_l^{\min}$ such that at least two arrivals are made per hour, for all lines $l \in L$. Gravitational decay parameter $\beta = 1.0$, maximum relative user cost increase $\varepsilon = 0.01$ and mode weights $\theta_1 = \theta_2 = \theta_3 = 1$ are also kept the same as in the Chicago case.

Seeing as the amount of pc4 communities is quite small ($|\tilde{V}_{\text{pc4}}^+| = 16$), and two of these regions are very sparsely populated with no bus service, taking roughly 10% of this would not suffice. Thus the community count is set higher at $\mathcal{K}_{\text{pc4}} = 4$, or 25%. The voronoi regions do not have these issues, so their community count can be set at $\mathcal{K}_{\text{vor}} = 10$ to be in line with the Chicago case. **These parameters may be finetuned after running, just like in the paper by Rumpf**

Chapter 5

Explanation of experiment

Chapter 6

Results

Chapter 7

Discussion

Appendix A

Pseudocode

The pseudocode for the solution algorithm discussed in section 3.7, directly taken from [RK21].

x represents the entirety of the flow vector. The vector $e_l \in \{0, 1\}^{|L|}$ represents a vector with a 1 in coordinate l and 0 elsewhere. Sets $\tilde{N}^+, \tilde{N}^-, \tilde{N}^\times$ represent the ADD, DROP and SWAP neighbors kept as candidates after the first pass of the neighborhood search, respectively, while $N^+, N^-, N^\times$ are the second-pass candidates. Bounds $\tilde{N}_{\max}, N_{\max}$ are the required numbers of candidates to be gathered during each pass. $Q_{\text{in}}^{\max}, Q_{\text{out}}^{\max}$ are the inner and outer nonimprovement counter bounds, respectively.

```

1: // INITIALIZATION
2:    $y^{(0)} \leftarrow y^*$  ▷ current real-world fleet sizes
3:    $y_{\text{best}} \leftarrow y^{(0)}$  ▷ best known solution and objective
4:    $\text{BestAccess} \leftarrow \text{Access}(y^{(0)})$  ▷ best known objective
5:    $\text{STM} \leftarrow \emptyset$  ▷ tabu move list, equipped with tabu tenures
6:    $\text{LTM} \leftarrow \emptyset$  ▷ attractive solution set
7:    $T \leftarrow T_0$  ▷ simulated annealing temperature
8:    $t \leftarrow t_0$  ▷ tabu tenure
9:    $Q_{\text{in}} \leftarrow 0$  ▷ inner nonimprovement counter
10:   $Q_{\text{out}} \leftarrow 0$  ▷ outer nonimprovement counter
11:  for  $k = 1, \dots, K$  do ▷ search for a predetermined number  $K$  of local moves
12:    // NEIGHBORHOOD SEARCH
13:     $\tilde{N}^+ \leftarrow \emptyset$  ▷ tentative ADD neighborhood
14:    repeat ▷ first pass (ADD)
15:       $l \leftarrow$  random element of  $L$  not yet considered during this loop
16:      if  $y^{(k)} + e_l$  satisfies vehicle limit bounds (3.1 - 3.3) then ▷ consider only design-feasible
17:        moves
18:        if  $e_l \notin \text{STM}$  or  $\text{Access}(y^{(k)} + e_l) > \text{BestAccess}$  then
19:           $\tilde{N}^+ \leftarrow \tilde{N}^+ \cup \{l\}$  ▷ keep non-tabu candidates (unless improved best)
20:        until  $|\tilde{N}^+| \geq \tilde{N}_{\max}$  or all lines in  $L$  have been considered
21:       $N^+ \leftarrow \emptyset$  ▷ final ADD neighborhood
22:      repeat
23:         $l \leftarrow$  unprocessed element of  $\tilde{N}^+$  with maximum  $\text{Access}(y^{(k)} + e_l)$ 
24:         $x^{(k)} \leftarrow \text{TransitAssignment}(y^{(k)} + e_l)$  ▷ calculate user flows for given move
25:        if  $(y^{(k)} + e_l, x^{(k)})$  satisfy operator and user cost constraints (3.5 - 3.6) then
26:           $N^+ \leftarrow N^+ \cup \{l\}$  ▷ keep only feasible candidates
27:        until  $|N^+| \geq N_{\max}$  or all lines in  $\tilde{N}^+$  have been processed
28:       $\tilde{N}^- \leftarrow \emptyset$  ▷ tentative DROP neighborhood
29:      repeat ▷ first pass (DROP)
30:         $l \leftarrow$  random element of  $L$  not yet considered during this loop
31:        if  $y^{(k)} - e_l$  satisfies vehicle limit bounds (3.1 - 3.3) then ▷ consider only design-feasible
32:          moves
33:          if  $-e_l \notin \text{STM}$  or  $\text{Access}(y^{(k)} - e_l) > \text{BestAccess}$  then
34:             $\tilde{N}^- \leftarrow \tilde{N}^- \cup \{l\}$  ▷ keep non-tabu candidates (unless improved best)
35:          until  $|\tilde{N}^-| \geq \tilde{N}_{\max}$  or all lines in  $L$  have been considered
36:         $N^- \leftarrow \emptyset$  ▷ final DROP neighborhood

```

```

35: repeat
36:    $l \leftarrow$  unprocessed element of  $\tilde{N}^-$  with maximum  $\text{Access}(y^{(k)} - e_l)$ 
37:    $x^{(k)} \leftarrow \text{TransitAssignment}(y^{(k)} - e_l)$   $\triangleright$  calculate user flows for given move
38:   if  $(y^{(k)} - e_l, x^{(k)})$  satisfy operator and user cost constraints (3.5 - 3.6) then
39:      $N^- \leftarrow N^- \cup \{l\}$   $\triangleright$  keep only feasible candidates
40:   until  $|N^-| \geq N_{\max}$  or all lines in  $\tilde{N}^-$  have been processed
41:   if  $|N^+| = |N^-| = 0$  then  $\triangleright$  handle the event of a failed search
42:     if bounds  $N_{\max}$  were used in the previous neighborhood search then
43:       restart neighborhood search while ignoring the bounds  $\tilde{N}_{\max}$ 
44:     else
45:        $Q_{\text{out}} \leftarrow Q_{\text{out}} + 1$   $\triangleright$  work towards diversification
46:       delete from STM the tabu rule with the shortest remaining tenure
47:       restart neighborhood search
48:    $\tilde{N}^\times \leftarrow \{(m, l) \in N^- \times N^+ : z_m = z_l\}$   $\triangleright$  set of all compatible DROP/ADD pairs
49:    $N^\times \leftarrow \emptyset$   $\triangleright$  final SWAP neighborhood
50:   repeat
51:      $(m, l) \leftarrow$  unprocessed element of  $N^\times$  with maximum  $\text{Access}(y^{(k)} - e_m) + \text{Access}(y^{(k)} + e_l)$ 
52:      $x^{(k)} \leftarrow \text{TransitAssignment}(y^{(k)} - e_m + e_l)$ 
53:     if  $y^{(k)} - e_m + e_l, x^{(k)}$  satisfy operator and user cost constraints (3.5 - 3.6) then
54:        $N^\times \leftarrow N^\times \cup \{(m, l)\}$   $\triangleright$  keep only feasible candidates
55:   until  $|N^\times| \geq N_{\max}$  or all moves in  $\tilde{N}^\times$  have been processed
56:    $\tilde{y}^{(k,1)} \leftarrow$  element of  $N^+ \cup N^- \cup N^{\text{times}}$  with greatest access  $\triangleright$  returning two best neighbors
57:    $\tilde{y}^{(k,2)} \leftarrow$  element of  $N^+ \cup N^- \cup N^{\text{times}}$  with second-greatest access
58:   // NEIGHBOR PROCESSING
59:   if  $\text{Access}(\tilde{y}^{(k,1)}) - \text{Access}(\tilde{y}^{(k)}) > \text{then}$ 
60:      $Q_{\text{out}} \leftarrow 0$   $\triangleright$  improvement iteration
61:      $t \leftarrow t_0$ 
62:      $y^{(k,1)} \leftarrow \tilde{y}^{(k,1)}$   $\triangleright$  move to improving neighbor
63:     add rules with tenure  $t$  to STM to prevent undoing the DROP/ADD moves of  $y^{(k,1)}$ 
64:     if  $\text{Access}(\tilde{y}^{(k,1)}) \geq \text{BestAccess}$  then
65:        $y_{\text{best}} \leftarrow \tilde{y}^{(k,1)}$   $\triangleright$  update best known solution
66:        $\text{BestAccess} \leftarrow \text{Access}(\tilde{y}^{(k,1)})$ 
67:   else
68:      $Q_{\text{out}} \leftarrow Q_{\text{out}} + 1$   $\triangleright$  nonimprovement iteration
69:      $Q_{\text{in}} \leftarrow Q_{\text{in}} + 1$ 
70:      $r \sim U[0, 1]$   $\triangleright$  random variable drawn uniformly from  $[0, 1]$ 
71:     if  $r < \exp\{-[\text{Access}(y^{(k)}) - \text{Access}(\tilde{y}^{(k,1)})]/T\}$  then
72:       increment  $t$   $\triangleright$  SA criterion passed
73:        $y^{(k+1)} \leftarrow \tilde{y}^{(k,1)}$   $\triangleright$  move to non-improving neighbor
74:       make undoing the DROP/ADD moves of  $y^{(k+1)}$  tabu with tenure  $t$ 
75:        $Q_{\text{in}} \leftarrow 0$ 
76:        $\text{LTM} \leftarrow \text{LTM} \cup \{\tilde{y}^{(k,2)}\}$   $\triangleright$  keep second-best neighbor as an attractive solution
77:     else
78:        $\text{LTM} \leftarrow \text{LTM} \cup \{\tilde{y}^{(k,1)}\}$   $\triangleright$  keep best neighbor as an attractive solution
79:   // UPKEEP
80:   if  $Q_{\text{in}} \geq Q_{\text{in}}^{\max}$  then
81:      $Q_{\text{in}} \leftarrow 0$   $\triangleright$  attempt to diversify
82:      $Q_{\text{out}} \leftarrow Q_{\text{out}} + 1$ 
83:      $y \leftarrow$  randomly chosen element of LTM
84:      $\text{LTM} \leftarrow \text{LTM} \setminus \{y\}$ 
85:      $y^{(k+1)} \leftarrow y$   $\triangleright$  move to a random attractive solution
86:   increment  $t$ 
87:   if  $Q_{\text{out}} \geq Q_{\text{out}}^{\max}$  then
88:      $t \leftarrow t_0$   $\triangleright$  attempt to intensify
89:   remove moves from STM whose tenures have elapsed
90:   apply cooling schedule to  $T$ 
91: return  $(y_{\text{best}}, \text{BestAccess})$ 

```

After the algorithm terminates and we obtain the best known solution, we conduct an exhaustive local search starting at y_{best} . In each iteration of this local search we consider every possible ADD, DROP, and SWAP move without considering tabu rules, always taking the best neighbor and terminating when no neighbors offer any improvement. The final result is guaranteed to be locally optimal.

Appendix B

Data used

All data used for this study. Community and travel data was created following methodology in Section 4. Other data, like bus-schedule, bus-stops and GP-office data, was only formatted after being aquired. All used .csv-files can be found on [Ock]. Note that data related to the bus-schedule and general practitioners may be outdated as these are prone to modification.

B.1 Communities

B.1.1 PC4

Population and location data for PC4 regions. Location is based on the rough centerpoint of each region, population is taken from [All25].

PC4	Latitude	Longitude	Population
2311	52,155929	4,489121	11.355
2312	52,16134	4,493255	15.260
2313	52,150299	4,501009	12.380
2314	52,150391	4,515009	6.200
2315	52,16334	4,510335	8.715
2316	52,168447	4,497329	10.720
2317	52,17676	4,510158	10.720
2318	52,178849	4,500912	3.250
2321	52,147943	4,481043	12.795
2322	52,139427	4,49774	225
2323	52,131764	4,480448	170
2324	52,144827	4,471378	9.970
2331	52,152021	4,449859	10.705
2332	52,160386	4,462739	10.545
2333	52,167158	4,467386	4.295
2334	52,174422	4,481533	2.805

B.1.2 Voronoi

Population and location data for generated Voronoi regions. Although regions are generated with bus stops as weightpoints, ID numbers do not match due to the voronoi generating algorithm used. Also note empty regions 88 and 92. Region 88 is an dummy region added by the voronoi algorithm and region 92 lies outside the study area, so no population is assigned.

ID	Latitude	Longitude	Population	ID	Latitude	Longitude	Population
0	52,1804	4,514149	3555	50	52,155501	4,454122	572
1	52,181422	4,510137	1472	51	52,173228	4,473882	1970
2	52,148038	4,481483	1660	52	52,147679	4,498162	1514
3	52,142884	4,478695	2417	53	52,15896	4,500553	2062
4	52,152654	4,443192	1215	54	52,158554	4,505135	1949
5	52,161507	4,485565	1507	55	52,154885	4,50135	2377
6	52,181015	4,502586	1159	56	52,162637	4,499515	1733
7	52,172564	4,486935	842	57	52,16131	4,49751	1702
8	52,173307	4,481789	1007	58	52,153439	4,473205	1295
9	52,178798	4,497163	809	59	52,157171	4,45736	918
10	52,163707	4,489994	2525	60	52,155844	4,494945	1207
11	52,168418	4,493743	3177	61	52,159378	4,497534	909
12	52,150492	4,50398	2417	62	52,156012	4,493392	1514
13	52,144599	4,503278	816	63	52,154928	4,49755	1064
14	52,14284	4,508426	1001	64	52,15094	4,494327	887
15	52,154821	4,507182	1858	65	52,149578	4,472554	724
16	52,163155	4,512309	2264	66	52,143317	4,476568	1623
17	52,152559	4,452693	699	67	52,147757	4,477632	1327
18	52,150612	4,44854	1070	68	52,150761	4,456047	615
19	52,165519	4,482503	1356	69	52,159128	4,463881	1173
20	52,170021	4,479999	379	70	52,157172	4,466626	896
21	52,169867	4,488101	1461	71	52,157464	4,459196	147
22	52,163192	4,486101	1059	72	52,170223	4,458579	1025
23	52,162264	4,484644	760	73	52,168172	4,457219	366
24	52,167636	4,505643	2329	74	52,163972	4,464921	367
25	52,175828	4,503727	1960	75	52,163113	4,465407	605
26	52,173831	4,498752	2588	76	52,161515	4,463423	845
27	52,175292	4,507802	2943	77	52,165125	4,460629	256
28	52,163696	4,506437	1652	78	52,166667	4,461172	314
29	52,166456	4,50647	710	79	52,169725	4,463896	508
30	52,167034	4,511677	3168	80	52,163458	4,469302	649
31	52,167034	4,511677	755	81	52,169368	4,471139	762
32	52,1519	4,514531	1041	82	52,162358	4,47352	811
33	52,148771	4,513102	2048	83	52,15968	4,470052	370
34	52,154994	4,512279	1076	84	52,163767	4,472055	1439
35	52,154554	4,517226	1300	85	52,158953	4,476015	464
36	52,15237	4,520981	955	86	52,16641	4,480336	955
37	52,146655	4,513888	2424	87	52,156909	4,473222	519
38	52,157532	4,491243	1173	88	52,170112	4,476591	0
39	52,150027	4,48891	1754	89	52,147558	4,449948	1763
40	52,153246	4,490145	613	90	52,146891	4,455354	965
41	52,149458	4,493246	1735	91	52,15494	4,462858	1950
42	52,147308	4,492968	1183	92	52,145792	4,460797	0
43	52,148271	4,486125	3072	93	52,144972	4,470851	416
44	52,145327	4,486637	1571	94	52,145487	4,471293	497
45	52,158179	4,48176	3210	95	52,14815	4,470362	826
46	52,154287	4,48292	1708	96	52,146619	4,467661	806
47	52,160053	4,481181	849	97	52,141466	4,470806	1716
48	52,155372	4,44578	1571	98	52,141399	4,473744	2406
49	52,155781	4,449805	1004				

B.2 Travel data

Bibliography

- [All25] Allecijfers. May 1, 2025. URL: <https://allecijfers.nl/postcode/>.
- [Møl94] Jesper Møller. *Lectures on Random Voronoi Tessellations*. Springer New York, NY, 1994. DOI: <https://doi.org/10.1007/978-1-4612-2652-9>.
- [Ned25] Patientenfederatie Nederland. May 2, 2025. URL: <https://www.zorgkaartnederland.nl/huisartsenpraktijk/leiden>.
- [Ock] Jeroen Ockers. URL: <https://github.com/Jerrie10/Bachelor-Thesis/tree/main/Data%20Used>.
- [Qbu25] Qbuzz. Feb. 3, 2025. URL: <https://zhn.qbuzz.nl/?direction=1&date=2025-02-10>.
- [RK21] Adam Rumpf and Hemanshu Kaul. “A public transit network optimization model for equitable access to social services”. eng. In: (2021).
- [SF89] Heinz Spiess and Michael Florian. “Optimal strategies: A new assignment model for transit networks”. In: *Transportation Research Part B: Methodological* 23.2 (1989), pp. 83–102. ISSN: 0191-2615. DOI: [https://doi.org/10.1016/0191-2615\(89\)90034-9](https://doi.org/10.1016/0191-2615(89)90034-9). URL: <https://www.sciencedirect.com/science/article/pii/0191261589900349>.
- [Swe25] Sweco. Sept. 18, 2025. URL: <https://www.sweco.nl/portfolio/het-gorlaeus-gebouw/>.
- [Wan22] M. Wang F.; Ye. “Optimization Method for Conventional Bus Stop Placement and the Bus Line Network Based on the Voronoi Diagram.” In: *Sustainability* 14.7918 (2022). DOI: <https://doi.org/10.3390/su14137918>.
- [Wei76] Jörgen W. Weibull. “An axiomatic approach to the measurement of accessibility”. In: *Regional Science and Urban Economics* 6.4 (1976), pp. 357–379. ISSN: 0166-0462. DOI: [https://doi.org/10.1016/0166-0462\(76\)90031-4](https://doi.org/10.1016/0166-0462(76)90031-4). URL: <https://www.sciencedirect.com/science/article/pii/0166046276900314>.